

# Event Log Architectures

## When quality matters

**Martin Thompson - @mjpt777**

**Event  
Sourcing**

**Command  
Sourcing**

**CQRS**



**Jimmy Bogard** 🍺 @jbogard · Jan 23

every Event Sourcing rescue project i've been on was also promised these things and found them not to be true (for their scenario)

1. more scalable
2. zero data loss
3. faster transactional performance
4. simpler system modeling
5. faster development timelines

💬 33

↻ 42

❤️ 223



[Show this thread](#)



**Jimmy Bogard** 🍺 @jbogard · Jan 23

every Event Sourcing rescue project i've been on was also promised these things and

1. I

2. 2

3. f

4. S

5. f



33

Show this t



**Kelly Sommers** @kellabyte · Jan 23

Replying to @jbogard

Let me correct this for you.

1. Requires more scale and complexity.

2. More touch points where data loss can occur.

3. Often less transactional guarantees that confuse end users and developers.

4. Much more complicated data flow.

5. Much longer development timelines.



2



3



64



**Jimmy Bogard** 🍺 @jbogard · Jan 23

Now hey that won't sell anything



1



40





**Jimmy Bogard** 🍺 @jbogard · Jan 23

every Event Sourcing rescue project i've been on was also promised these things and

1. I

2. 2

3. f

4. S

5. f

33

Show this t



**Kelly Sommers** @kellabyte · Jan 23

Replying to @jbogard

Let me correct this for you.



**Iforite** @lforite · Jan 25

Replying to @jbogard

Event sourcing is very hard. On the surface it's simple and solve all your problems, but deep down it's a lot of things to understand and think through. The problem with es/cqrd is that no one understand the same



5. Much longer development timelines.

2

3

64



**Jimmy Bogard** 🍺 @jbogard · Jan 23

Now hey that won't sell anything



1

40



# **Fowler vs Event Store**

# The Dark Side of Event Sourcing: Managing Data Conversion

Michiel Overeem<sup>1</sup>, Marten Spoor<sup>1</sup>, and Slinger Jansen<sup>2</sup>

<sup>1</sup>{m.overeem, m.spoor}@afas.nl, Department Architecture and Innovation, AFAS Software, The Netherlands

<sup>2</sup>slinger.jansen@uu.nl, Department of Information and Computing Sciences, Utrecht University, The Netherlands

**Abstract**—Evolving software systems includes data schema changes, and because of those schema changes data has to be converted. Converting data between two different schemas while continuing the operation of the system is a challenge when that system is expected to be available always. Data conversion in event sourced systems introduces new challenges, because of the relative novelty of the event sourcing architectural pattern, because of the lack of standardized tools for data conversion, and because of the large amount of data that is stored in typical event stores. This paper addresses the challenge of schema evolution and the resulting data conversion for event sourced systems. First

upgrade strategies that are fast, efficient, and seamless have to be designed and implemented.

One of the architectural patterns that in recent years emerged in the development of cloud systems is Command Query Responsibility Segregation (CQRS). The pattern was introduced by Young [5] and Dahan [6], and the goal of the pattern is to handle actions that change data (those are called commands) in different parts in the system than requests that ask for data (called queries). By separating the command-side (the part that

# Events, Logs, & State Machines



# State Machines

**Input  $\times$  State  $\rightarrow$  State**

# State Machines

**Input  $\times$  State  $\rightarrow$  State**

**Input  $\times$  State  $\rightarrow$  Output**

# Replicated State Machines

**Ordered Inputs**  
**+**  
**Deterministic Execution**  
**=>**  
**Same State & Outputs**

# Event Log

## (Ordered Inputs)

**Inputs vs Changes vs Outputs**

***Don't throw away data  
without good reason!***

# Domain Models



**Jessica Joy Kerr**

@jessitron



If these promises are false (and they are; I agree with @jbogard), then why WOULD we use event sourcing?

IMO, it's as close as we can get to modeling the real world.

[blog.jessitron.com/2020/01/24/cap...](https://blog.jessitron.com/2020/01/24/cap...)



**Jimmy Bogard**



@jbogard · Jan 23

every Event Sourcing rescue project i've been on was also promised these things and found them not to be true (for their scenario)

[Show this thread](#)

1. more scalable



**Modelling free from the  
restrictions of data stores**

# **Development Considerations**

# Development Considerations

- 1. Choose only deterministic algorithms**

## Development Considerations

- 1. Choose only deterministic algorithms**
- 2. Validate Inputs before mutation**

## Development Considerations

- 1. Choose only deterministic algorithms**
- 2. Validate Inputs before mutation**
- 3. Avoid Head-of-Line blocking**

## Development Considerations

- 1. Choose only deterministic algorithms**
- 2. Validate Inputs before mutation**
- 3. Avoid Head-of-Line blocking**
- 4. Version messages & extend via options**

# Snapshots

# Snapshots

- 1. Separate model from implementation**



# Snapshots

- 1. Separate model from implementation**
- 2. Choose an efficient & compact codec**

# Snapshots

- 1. Separate model from implementation**
- 2. Choose an efficient & compact codec**
- 3. Version everything**

# Snapshots

- 1. Separate model from implementation**
- 2. Choose an efficient & compact codec**
- 3. Version everything**
- 4. Store in a form that is easy to distribute**

# **Time & Timers**

## Time & Timers

1. **Don't read the system clock**

## Time & Timers

1. **Don't read the system clock**
2. **Don't read the system clock!**

## Time & Timers

1. **Don't read the system clock**
2. **Don't read the system clock!**
3. **Timestamp events / messages**

## Time & Timers

1. **Don't read the system clock**
2. **Don't read the system clock!**
3. **Timestamp events / messages**
4. **Manage timers centrally and reliably**



# Integration

# Integration

- 1. Adapt other systems with gateways**

# Integration

- 1. Adapt other systems with gateways**
- 2. Treat RDBMSs like a ledger**

# Integration

- 1. Adapt other systems with gateways**
- 2. Treat RDBMSs like a ledger**
- 3. Negotiate shared state on start-up**

# Integration

- 1. Adapt other systems with gateways**
- 2. Treat RDBMSs like a ledger**
- 3. Negotiate shared state on start-up**
- 4. Track responses with timers**

# Quality Attributes

***Non-Functional Requirements?***

# nonfunctional

adjective \ nän-'fən(k)-shnəl \

**: having no function : serving or performing  
no useful purpose**



# nonfunctional

adjective \ nän-'fən(ɡ)-shnəl \

: having no function : **serving or performing  
no useful purpose**

: not performing or able to perform a regular  
function

***Words Matter***

# Fault Tolerance

# Fault Tolerance

***Primary / Secondary***  
**vs**  
***Consensus***

***Leslie Lamport - Paxos***

***Barbara Liskov - Viewstamp Replication***

***Ken Birman - Virtual Synchrony***

## In Search of an Understandable Consensus Algorithm (Extended Version)

Diego Ongaro and John Ousterhout  
Stanford University

### Abstract

Raft is a consensus algorithm for managing a replicated log. It produces a result equivalent to (multi-)Paxos, and it is as efficient as Paxos, but its structure is different from Paxos; this makes Raft more understandable than Paxos and also provides a better foundation for building practical systems. In order to enhance understandability, Raft separates the key elements of consensus, such as leader election, log replication, and safety, and it enforces a stronger degree of coherency to reduce the number of states that must be considered. Results from a user study demonstrate that Raft is easier for students to learn than

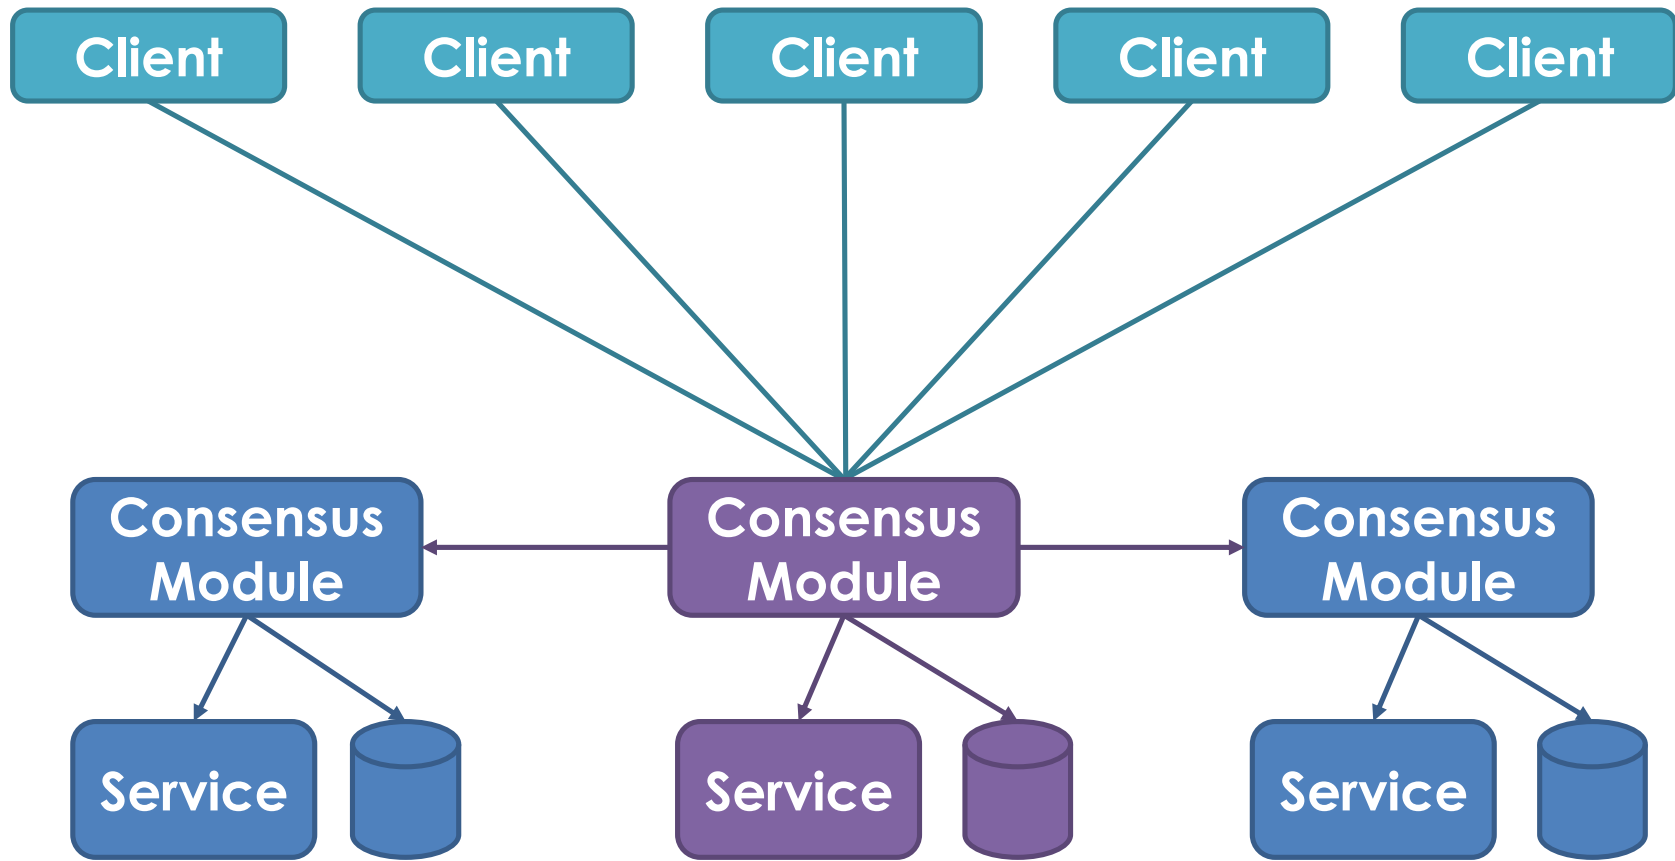
state space reduction (relative to Paxos, Raft reduces the degree of nondeterminism and the ways servers can be inconsistent with each other). A user study with 43 students at two universities shows that Raft is significantly easier to understand than Paxos: after learning both algorithms, 33 of these students were able to answer questions about Raft better than questions about Paxos.

Raft is similar in many ways to existing consensus algorithms (most notably, Oki and Liskov's Viewstamped Replication [29, 22]), but it has several novel features:

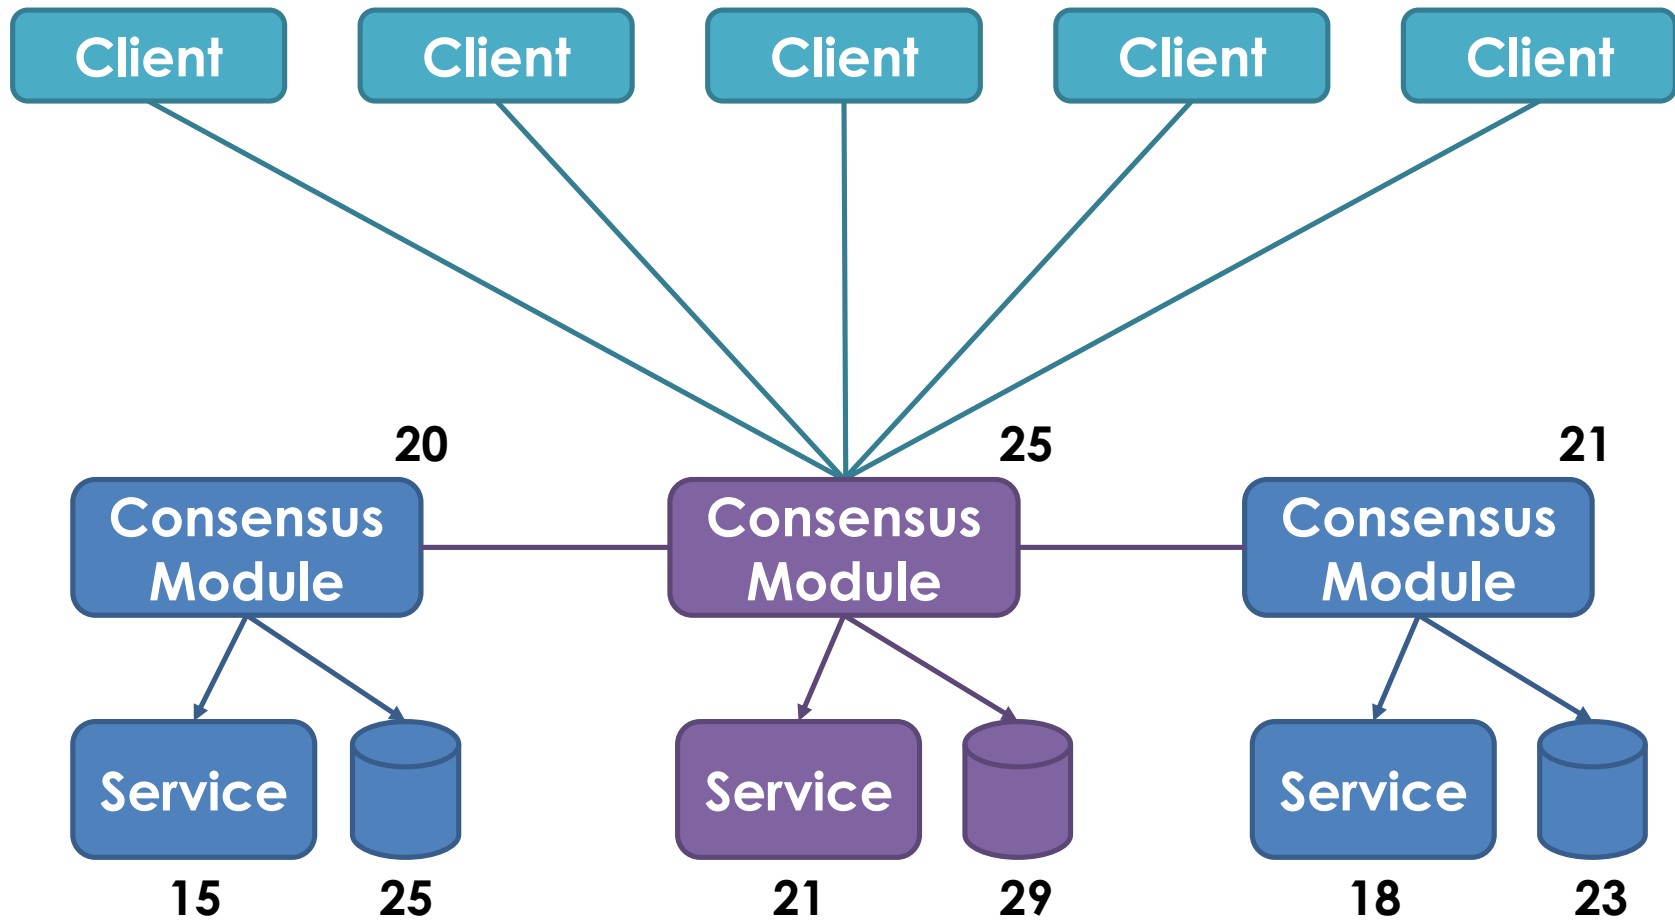
- **Strong leader:** Raft uses a stronger form of leadership than other consensus algorithms. For example,

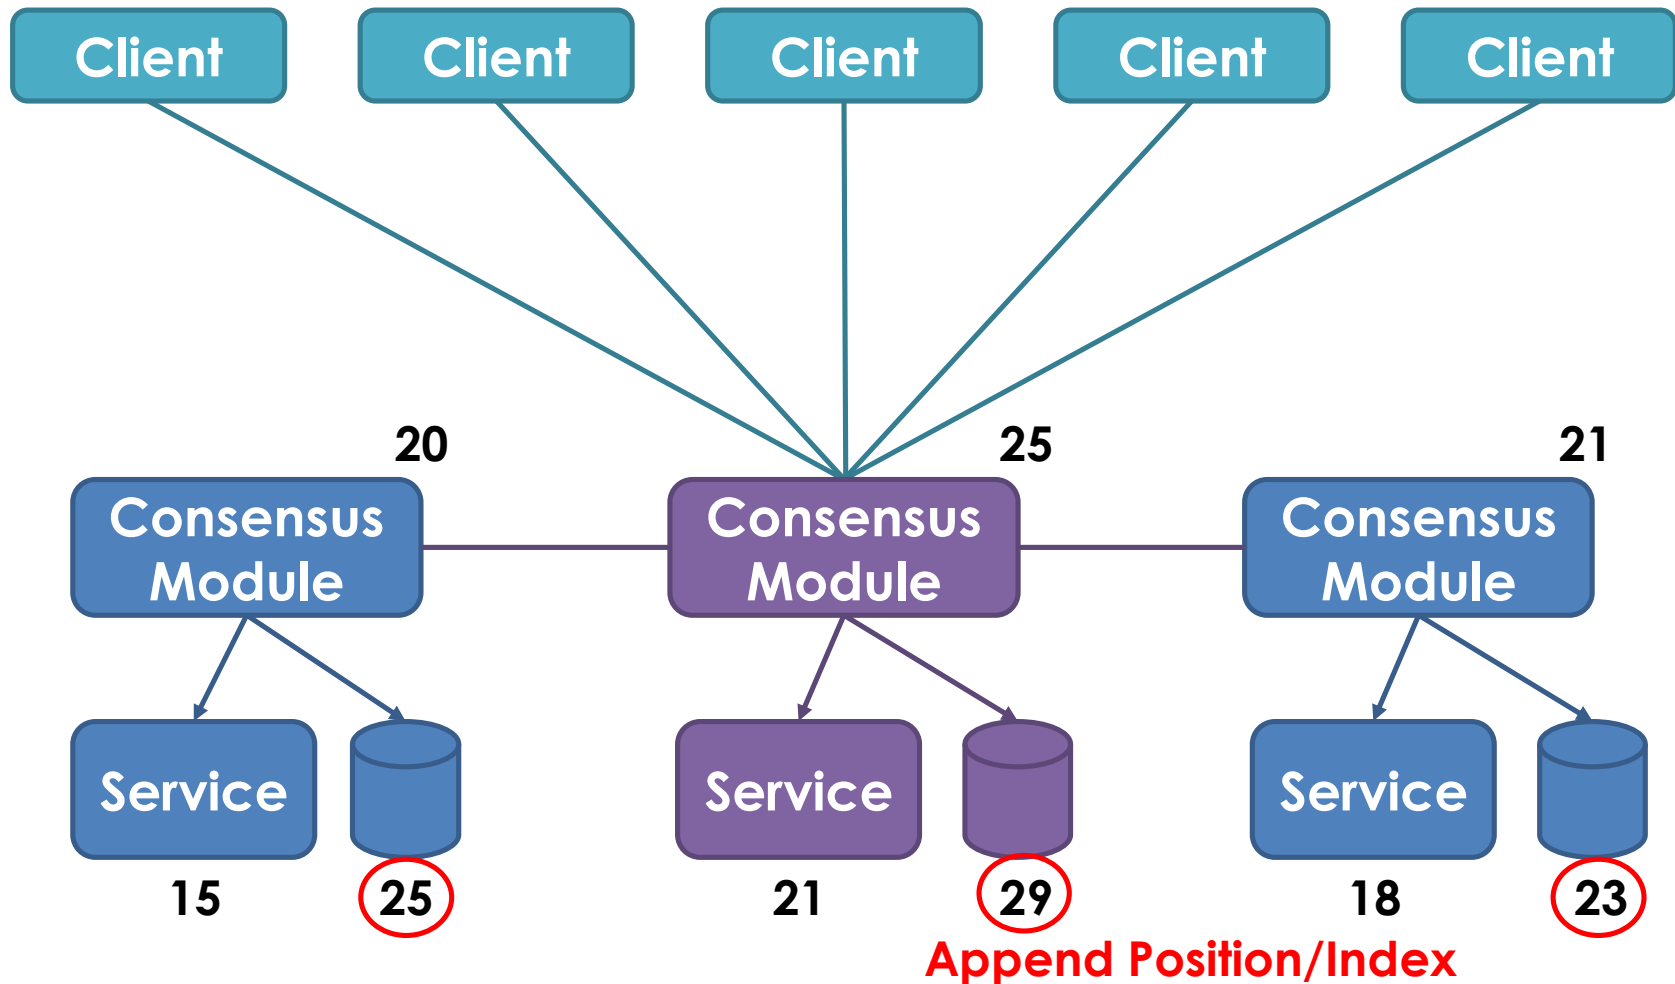
# Raft Safety Guarantees

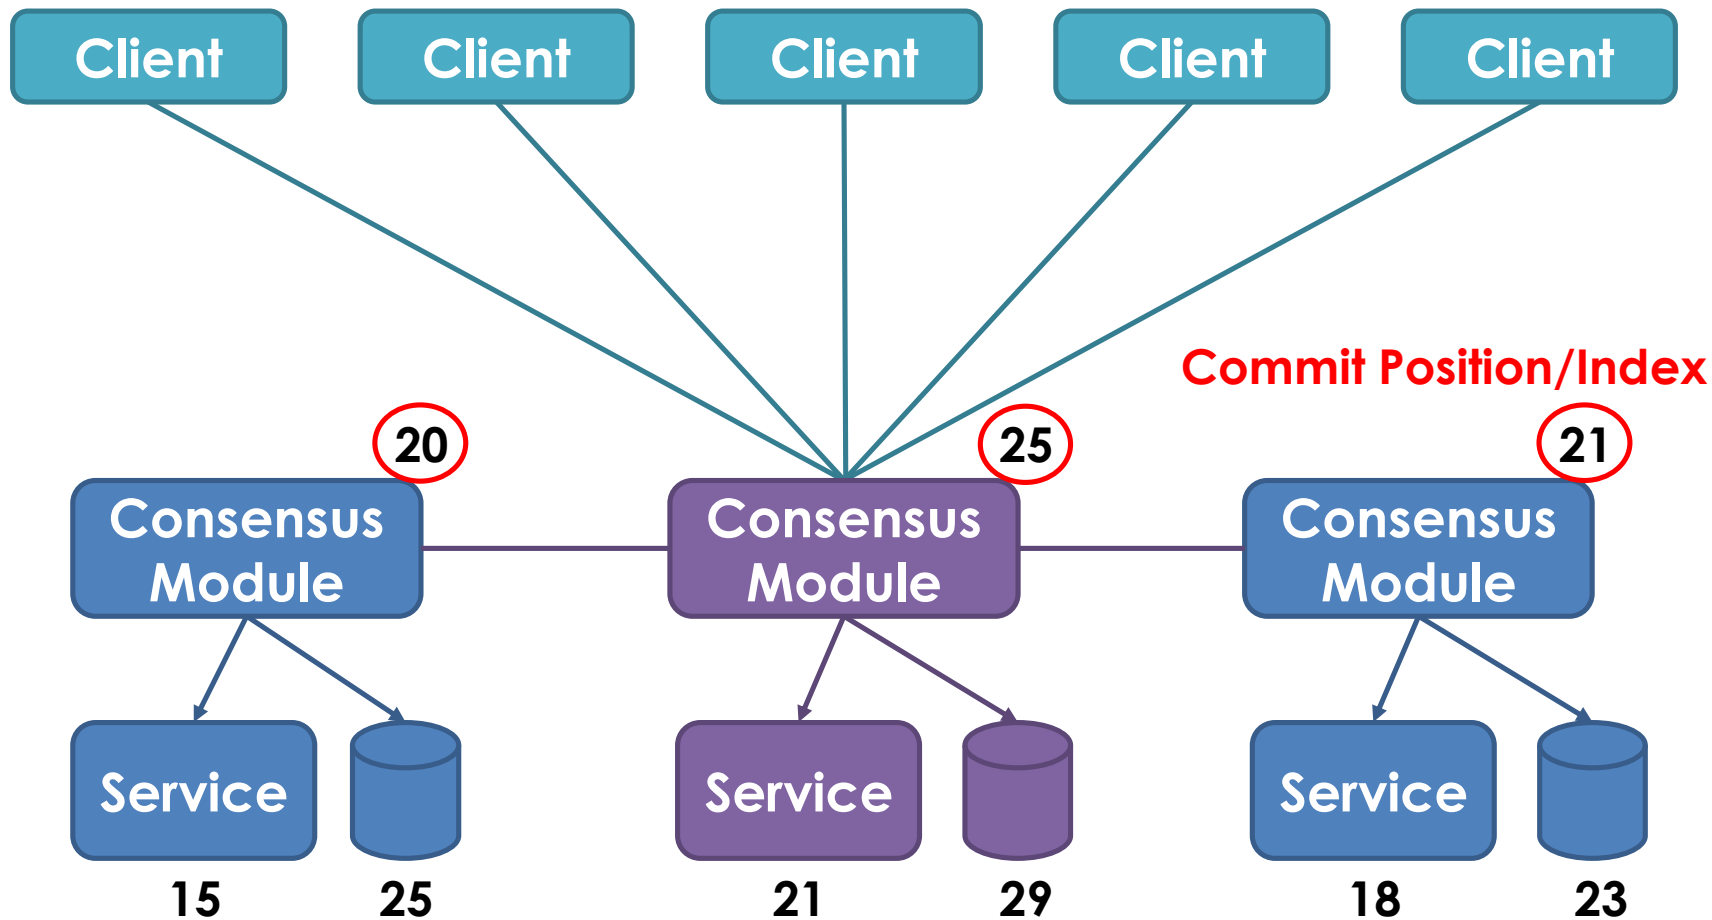
- **Election Safety**
- **Leader Append-Only**
- **Log Matching**
- **Leader Completeness**
- **State Machine Safety**

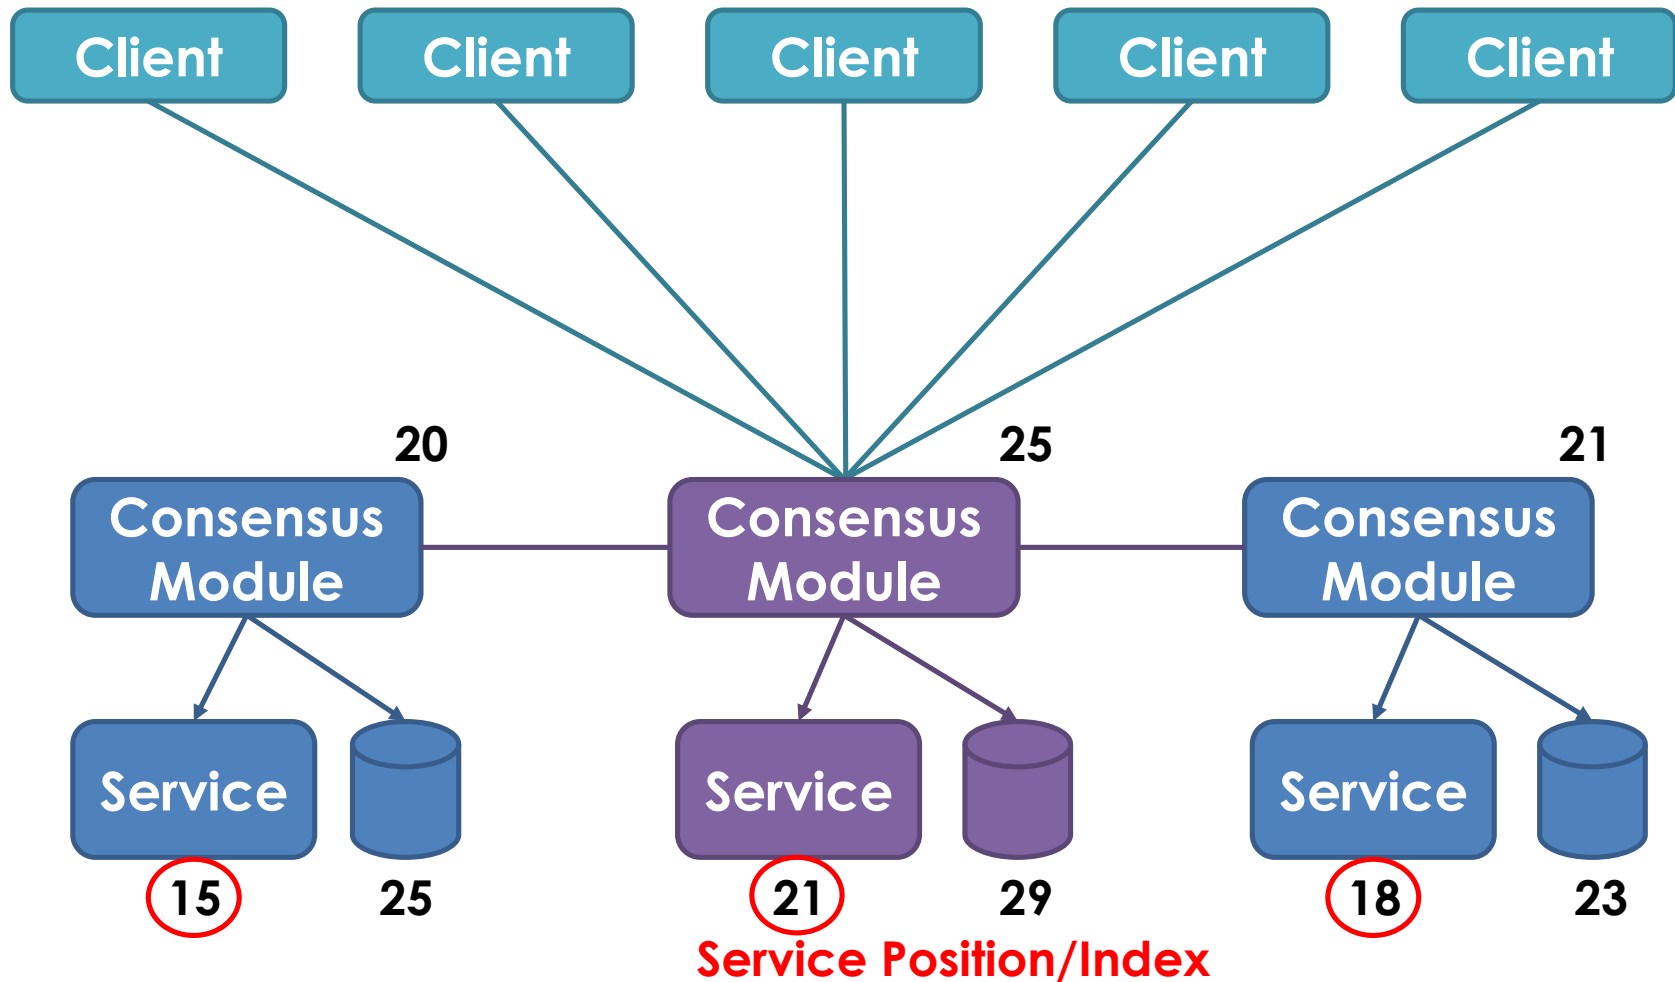


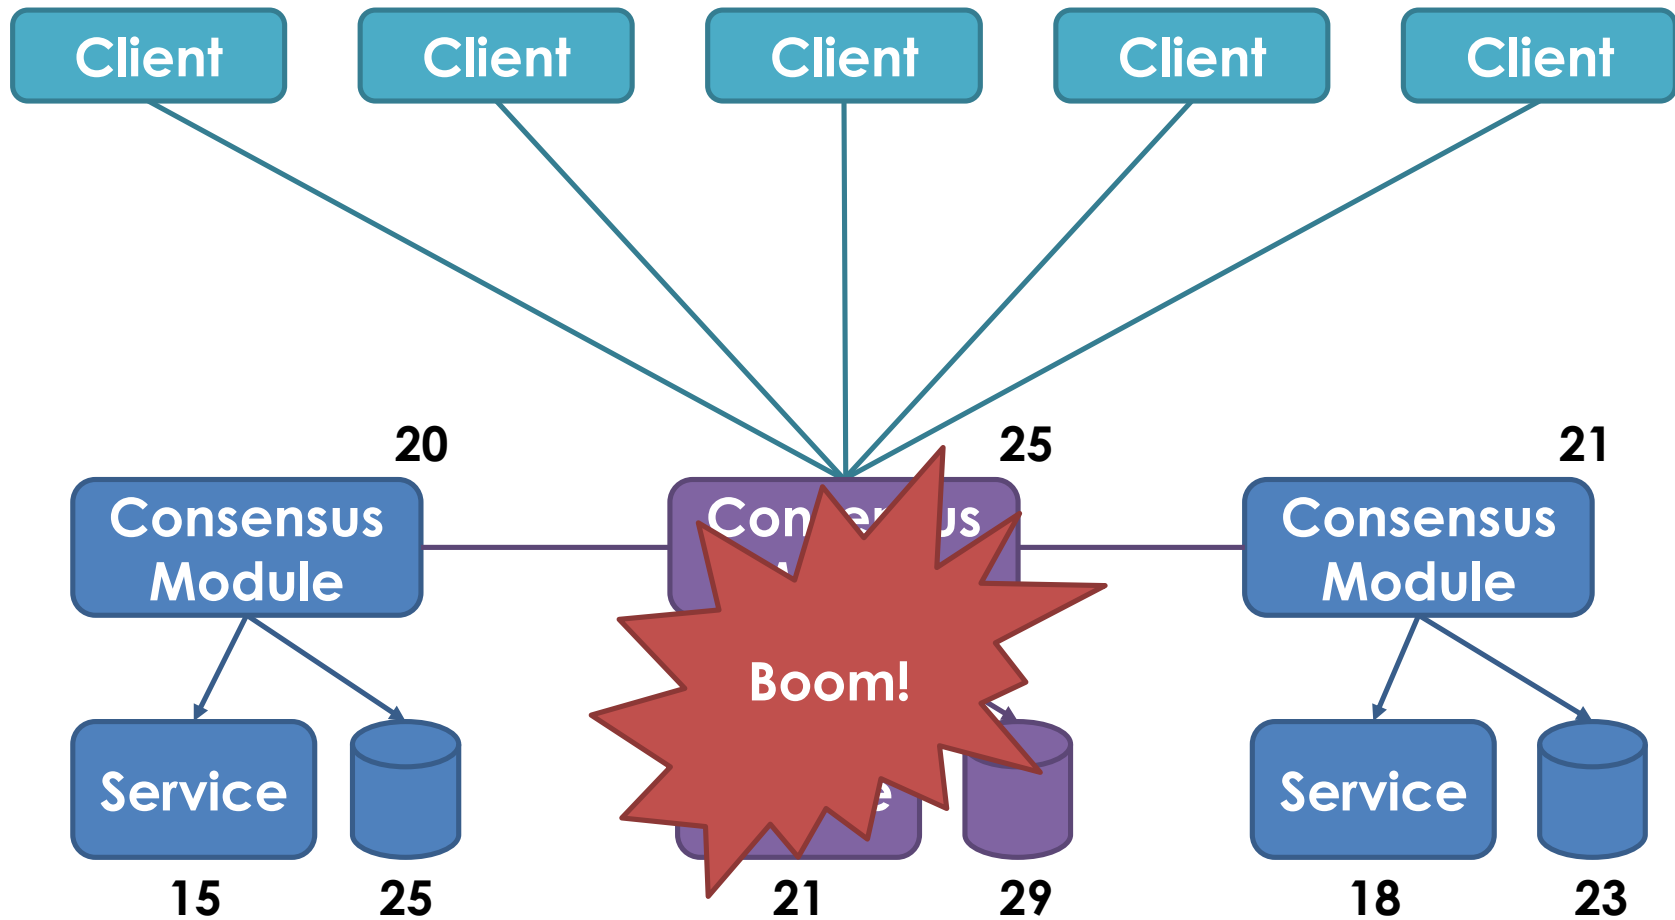


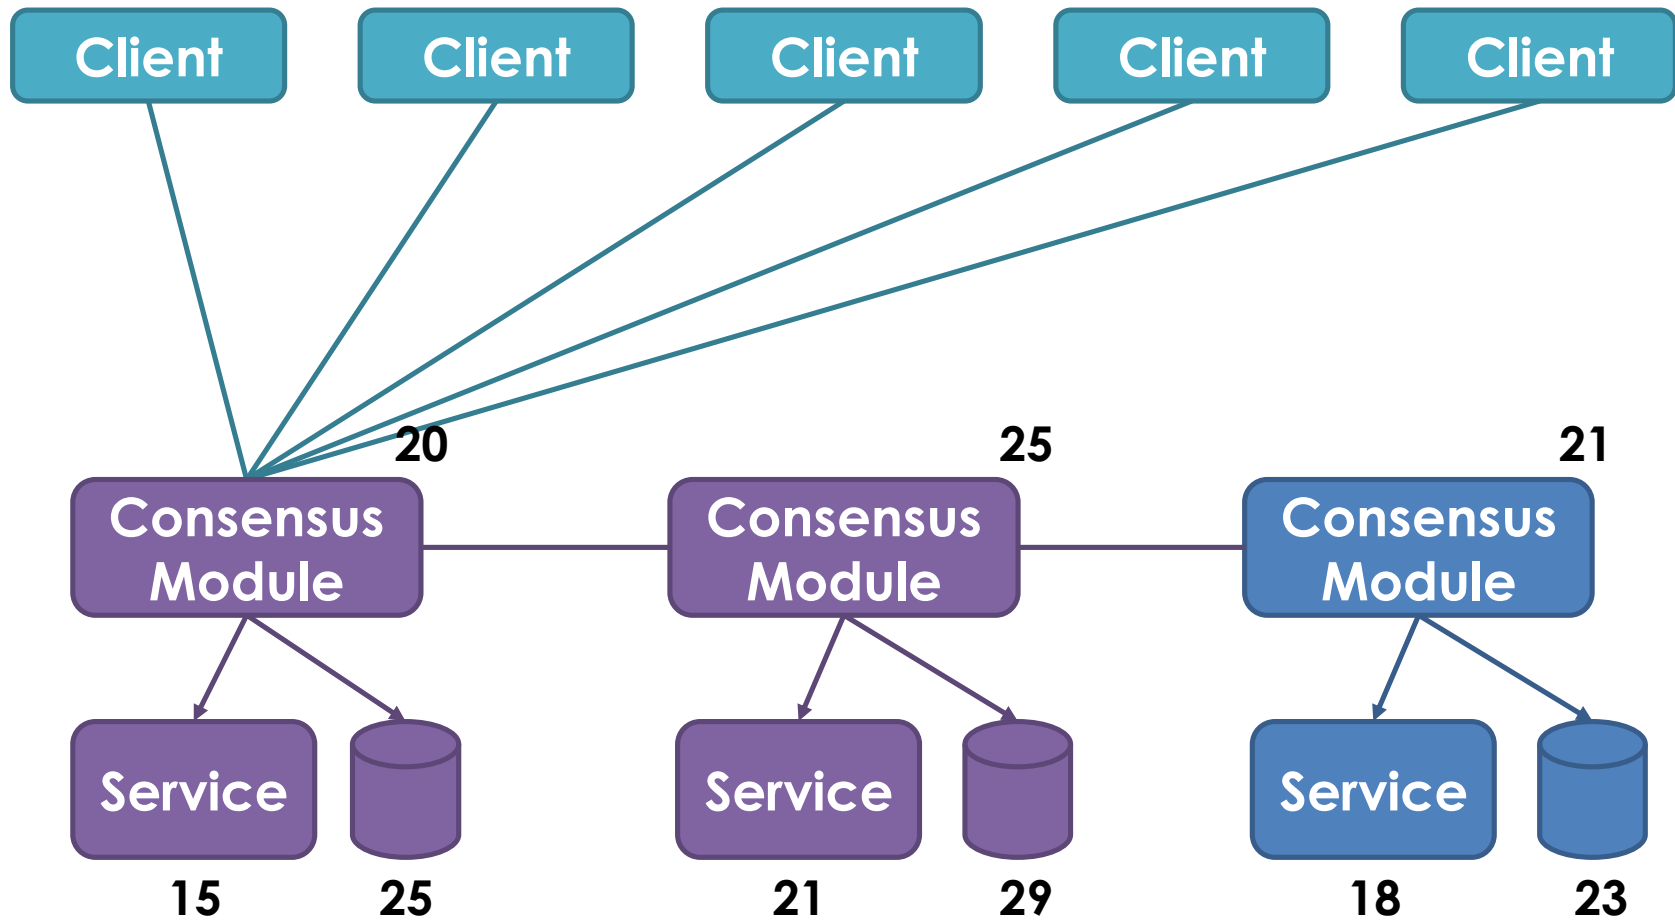












**Performance**

# Performance

- 1. Don't pollute the business logic thread**



## Performance

- 1. Don't pollute the business logic thread**
- 2. Choose appropriate data structures**

## Performance

- 1. Don't pollute the business logic thread**
- 2. Choose appropriate data structures**
- 3. Batch to amortise expensive costs**

# Performance

- 1. Don't pollute the business logic thread**
- 2. Choose appropriate data structures**
- 3. Batch to amortise expensive costs**
- 4. Continuously Perf Test and Profile**

# Scalability

# Scalability

## 1. Build multiple models as Services

# Scalability

- 1. Build multiple models as Services**
- 2. Shard model into Services**

# Scalability

- 1. Build multiple models as Services**
- 2. Shard model into Services**
- 3. Message between Services via the log**

# Scalability

- 1. Build multiple models as Services**
- 2. Shard model into Services**
- 3. Message between Services via the log**
- 4. Build replicas for queries, with caution**



**Wrapping up...**

# **Model Fidelity**

**10+ years ago...**

**@mjpt777**

**<https://github.com/real-logic/aeron>**

***“The future is already here – it's just not evenly distributed”***

**- William Gibson**